

PASTA Algorithm: Predictive Adaptive Scheduling with Timed Aging

Document Date: May 26, 2026

Subject: Deep analysis of PASTA scheduling algorithm for Apache Spark

Authors: Adaptive Scheduling Research

Audience: Students, systems engineers, algorithm enthusiasts

Table of Contents

1. [Executive Summary](#)
 2. [Problem Statement](#)
 3. [PASTA Algorithm Overview](#)
 4. [Core Components](#)
 5. [Mathematical Foundation](#)
 6. [Algorithm Walkthrough](#)
 7. [Comparison with Alternatives](#)
 8. [Implementation Details](#)
 9. [Complexity Analysis](#)
 10. [Advantages & Limitations](#)
 11. [Real-World Examples](#)
 12. [Empirical Results](#)
-

Executive Summary

What is PASTA?

PASTA (Predictive Adaptive Scheduling with Timed Aging) is a **hybrid scheduling algorithm** for Apache Spark that combines three key ideas:

1. **Shortest Job First (SJF)** — Schedule shorter predicted jobs first to minimize average latency
2. **Aging** — Prevent starvation by gradually increasing priority of long-waiting jobs
3. **Dynamic Tiers** — Adapt resource allocation thresholds to current workload mix
4. **Feedback Loop** — Continuously retrain ML model on actual execution times

Why PASTA?

Pure SJF Problem: A 60-second job can starve indefinitely if 1-second jobs keep arriving.

Pure FIFO Problem: A 1-minute job at the head blocks all 5-second jobs behind it.

PASTA Solution: Use SJF (minimize latency) + Aging (prevent starvation) + ML prediction (accurate timing) + Dynamic tier boundaries (adapt to workload).

Result

- ✓ **25-30% speedup** on heterogeneous workloads
- ✓ **No job starvation** (aging guarantee)
- ✓ **Adaptive to workload** (dynamic tiers)
- ✓ **Self-improving** (feedback retraining)

Problem Statement

The Scheduling Challenge in Apache Spark

When multiple jobs are queued for execution:

Job Queue (in order of arrival):

Job 1: FIFO default (60s)	← Head of queue, blocks all ot
Job 2: Interactive (2s)	
Job 3: Analytics (5s)	
Job 4: Batch (90s)	

Spark's Default FIFO Scheduler:

- Job 1 runs for 60s (all 2, 3, 4 wait)
- Job 2 runs for 2s
- Job 3 runs for 5s
- Job 4 runs for 90s
- **Total time: 157s**
- **Average latency: $(60 + 62 + 67 + 157) / 4 = 86.5s$**
- **Job 2 waited 60s for just 2s of work** ❌

Optimal SJF Scheduler:

- Job 2 runs for 2s (fastest)
- Job 3 runs for 5s
- Job 1 runs for 60s
- Job 4 runs for 90s

- **Total time: 157s** (same makespan)
- **Average latency: $(2 + 7 + 67 + 157) / 4 = 58.25s$** ✓
- **Job 2 completes in 2s** ✓

But there's a catch: How do we know job durations in advance?

The Starvation Problem

With SJF + Stream of Short Jobs:

Queue at time 0:

Job A: Large (120s)	← SJF puts this last
Job B: Small (2s)	

Execution:

[0-2s] Job B finishes
 New Job C arrives (2s) ← SJF adds to front
 [2-4s] Job C finishes
 New Job D arrives (2s) ← SJF adds to front
 [4-6s] Job D finishes
 New Job E arrives (2s) ← SJF adds to front
 ...
 [100s] Job A has NOT STARTED YET! ❌ STARVED

Pure SJF Allows Starvation: Long jobs never get scheduled if short jobs keep arriving.

The Prediction Challenge

Without ML Model:

- How do we know a job will take 60s vs 90s vs 2s?
- Cannot use SJF without predictions
- Must fall back to FIFO (no benefit)

With Inaccurate Predictions:

- Predict 2s, actual 60s → Wrong priority order → Worse than FIFO
- Predict 60s, actual 2s → Wastes resources on overprovisioning

PASTA's Answer: Train ML model on job features (input size, partition count, etc.)

The Resource Allocation Challenge

Fixed Tier Boundaries (Adaptive Scheduler's Approach):

Small threshold: 10s (fixed)

Large threshold: 60s (fixed)

Workload 1 (typical):

Jobs: 2s, 5s, 8s, 15s, 45s, 120s

✓ Tiers work well

Workload 2 (light):

Jobs: 0.5s, 0.8s, 1.2s

✗ All classified as "small" (threshold too high)

Resources underutilized

Workload 3 (heavy):

Jobs: 30s, 35s, 40s, 90s, 120s, 180s

✗ Many "medium" jobs (threshold too low)

Some resources underallocated

PASTA's Answer: Compute thresholds from current queue distribution (percentiles)

PASTA Algorithm Overview

Three Key Innovations

1 Predictive Priority (PASTA Score)

Combines **SJF** (short jobs first) with **Aging** (long-waiting jobs boosted):

$$\text{priority}_i = \alpha \times \text{norm_sjf}_i + \beta \times \text{norm_aging}_i$$

Where:

$$\text{norm_sjf}_i = 1 - (\text{predicted}_i - \text{min_predicted}) / (\text{max_pre}$$

$$\text{norm_aging}_i = \text{wait_time}_i / (\text{max_wait_time} + \epsilon)$$

$$\alpha = 0.70 \text{ (SJF weight)}$$

$$\beta = 0.30 \text{ (Aging weight)}$$

Interpretation:

- $\text{norm_sjf}_i = 1.0 \rightarrow$ shortest job $\rightarrow$ gets 70% priority boost
 - $\text{norm_aging}_i = 1.0 \rightarrow$ waited longest $\rightarrow$ gets 30% priority boost
 - **Balance:** SJF prevents long delays, Aging prevents starvation
-

2 Dynamic Tier Boundaries (Percentile-Based)

Instead of fixed thresholds (10s / 60s), compute from current queue:

```
small_threshold = 33rd percentile of predicted times  
medium_threshold = 66th percentile of predicted times
```

Example:

Predicted times: [0.5s, 1.2s, 2.0s, 5.5s, 8.0s, 15.0s, 45.0s]

33rd percentile $\approx$ 2.5s

66th percentile $\approx$ 30.0s

Result:


small = < 2.5s [0.5, 1.2, 2.0]

medium = 2.5-30s [5.5, 8.0, 15.0, 45.0]

large = > 30s [120.0]

Benefit: Automatically adapts to workload

3 Feedback Loop (Continuous Learning)



Benefit: Model continuously improves over time

Core Components

Component 1: Execution Time Predictor

Input: Job specification

```
{  
    "job_type": "small|medium|large",  
    "input_size_mb": 256,  
    "num_partitions": 8,  
    # num_stages, num_tasks, peak_memory_mb, cpu_seconds infer  
}
```

Output: Predicted execution time (seconds)

predicted_time = 8.3 # seconds

Model: Random Forest + XGBoost ensemble

- **Random Forest:** Robust, fast inference, handles non-linearity
- **XGBoost:** Higher accuracy, handles regularization

Features Used:

Feature	Type	Example	Impact
job_type	Categorical	"small"	Job determines parallelism
input_size_mb	Continuous	256	Larger input → longer execution
num_partitions	Continuous	8	Parallelism affects duration
num_stages	Continuous	3	More stages → more overhead
num_tasks	Continuous	24	Total work units
peak_memory_mb	Continuous	512	Memory pressure affects speed
cpu_seconds	Continuous	50	Computational work
hour_of_day	Categorical	14	Temporal pattern (day vs night)

Accuracy: ~88-92% R² (correlation with actual times)

Component 2: Priority Computation Engine

Algorithm:

Input:

- All predicted times: $[t_1, t_2, \dots, t_n]$
- All wait times: $[w_1, w_2, \dots, w_n]$
- Current job i

Process:

1. Compute $\text{min_predicted} = \min(\text{predictions})$
2. Compute $\text{max_predicted} = \max(\text{predictions})$
3. Compute $\text{max_wait} = \max(\text{wait_times})$
4. $\text{norm_sjf_i} = 1 - (\text{predicted_i} - \text{min_predicted}) / (\text{max_pre} - \text{min_predicted})$
→ Normalize to $[0, 1]$, where 1 = shortest, 0 = longest
5. $\text{norm_aging_i} = \text{wait_i} / (\text{max_wait} + \epsilon)$
→ Normalize to $[0, 1]$, where 0 = just arrived, 1 = waitec
6. $\text{priority_i} = 0.70 \times \text{norm_sjf_i} + 0.30 \times \text{norm_aging_i}$

Output:

- $\text{priority_i} \in [0, 1]$ (higher = schedule sooner)

Example:

Queue (current_time = 100s):

Job A: predicted=60s, arrived_at=0s (wait=100s)
Job B: predicted=2s, arrived_at=95s (wait=5s)
Job C: predicted=30s, arrived_at=98s (wait=2s)

Normalization values:

min_predicted = 2s
max_predicted = 60s
max_wait = 100s

Job A:

$\text{norm_sjf} = 1 - (60 - 2) / (60 - 2 + \epsilon) = 1 - 1.0 = 0.0$ (lc
 $\text{norm_aging} = 100 / 100 = 1.0$ (w
 $\text{priority_A} = 0.70 \times 0.0 + 0.30 \times 1.0 = 0.30$

Job B:

$\text{norm_sjf} = 1 - (2 - 2) / 58 = 1.0$ (s
 $\text{norm_aging} = 5 / 100 = 0.05$ (j
 $\text{priority_B} = 0.70 \times 1.0 + 0.30 \times 0.05 = 0.715$

Job C:

$\text{norm_sjf} = 1 - (30 - 2) / 58 = 1 - 0.483 = 0.517$ (me
 $\text{norm_aging} = 2 / 100 = 0.02$ (j
 $\text{priority_C} = 0.70 \times 0.517 + 0.30 \times 0.02 = 0.368$

Ordering (by priority, descending):

1. Job B (0.715) → Schedule first ✓
2. Job C (0.368)
3. Job A (0.30) → This would starve with pure SJF, but needs

Key Insight:

- **Batch mode (all arrive at same time):** $w_i = 0$ for all → PASTA reduces to pure SJF
- **Streaming mode (arrive over time):** w_i varies → Aging prevents starvation

Component 3: Dynamic Tier Boundary Calculator

Algorithm:

```

def compute_dynamic_tiers(predictions: list[float]) -> tuple[float, float]:
    """
    Return (low_threshold, high_threshold) based on 33rd / 66th
    of current queue's predicted times.
    """
    low = np.percentile(predictions, 33)  # 33rd percentile
    high = np.percentile(predictions, 66)  # 66th percentile

    # Ensure gap between thresholds
    if high <= low:
        high = low * 2 + ε

    return low, high

def get_tier_dynamic(predicted: float, low: float, high: float) -> str:
    if predicted <= low:
        return "small"
    elif predicted <= high:
        return "medium"
    else:
        return "large"

```

Example:

Workload 1 (mixed sizes):

Predictions: [1s, 2s, 3s, 8s, 15s, 45s, 90s, 180s]

p33 = 3.67s

p66 = 65.0s

small	< 3.67s	[1, 2, 3]
medium	3.67-65s	[8, 15, 45, 90]
large	> 65s	[180]

Workload 2 (light, all small):

Predictions: [0.5s, 0.8s, 1.0s, 1.2s]

p33 = 0.83s

p66 = 1.1s

small	< 0.83s	[0.5]
medium	0.83-1.1s	[0.8, 1.0]
large	> 1.1s	[1.2]

→ Adaptively scales thresholds down!

Component 4: Resource Allocator

Mapping: Predicted time → Spark configuration

```

def get_config_dynamic(predicted: float, low: float, high: float):
    tier = get_tier_dynamic(predicted, low, high)

    if tier == "small":
        return {
            "spark.executor.memory": "512m",
            "spark.executor.cores": "1",
            "spark.sql.shuffle.partitions": "4",
            "spark.driver.memory": "1g",
        }
    elif tier == "medium":
        return {
            "spark.executor.memory": "1g",
            "spark.executor.cores": "2",
            "spark.sql.shuffle.partitions": "8",
            "spark.driver.memory": "1g",
        }
    else: # large
        return {
            "spark.executor.memory": "2g",
            "spark.executor.cores": "2",
            "spark.sql.shuffle.partitions": "16",
            "spark.driver.memory": "2g",
        }

```

Rationale:

- **Small jobs:** Minimal overhead, quick turnaround
 - **Medium jobs:** Balanced parallelism
 - **Large jobs:** Maximize parallelism with more partitions
-

Component 5: Feedback Retrainer

Algorithm:

After each benchmark run:

1. Collect actual execution times:
 - └ For each completed job: `actual_time_i`
2. Compute prediction error:
 - └ $\text{error}_i = | \text{predicted}_i - \text{actual}_i | / \text{actual}_i$
3. Append to training CSV:
 - └ `job_type, input_size_mb, num_partitions, ...`
 - └ `num_stages, num_tasks, peak_memory_mb, cpu_seconds`
 - └ `execution_time_sec, hour_of_day`
4. Retrain ML models:
 - └ `RandomForest.fit(X_train, y_train)`
 - └ `XGBoost.fit(X_train, y_train)`
5. Persist models:
 - └ `rf_model.joblib`
 - └ `xgb_model.joblib`
6. Next predictions use retrained models
 - └ Calibrated to local hardware!

Why Feedback?

Initial scenario:

Predictions:	[2s, 5s, 8s, 60s, 120s]
Actual:	[2.1s, 5.2s, 7.9s, 62s, 118s]
Error:	[5%, 4%, 1%, 3%, 1%] ← Good!

After hardware change (added more RAM):

Without feedback:

Predictions:	[2s, 5s, 8s, 60s, 120s] (unchanged)
Actual:	[1.5s, 3.8s, 5.9s, 45s, 85s]
Error:	[33%, 32%, 26%, 25%, 29%] × Very bad!

With feedback (after 1 run):

Retrained on new actual times	
Predictions:	[1.5s, 3.8s, 5.9s, 45s, 85s]
Error:	[0%, 0%, 0%, 0%, 0%] ✓ Perfect!

Mathematical Foundation

Part 1: Queueing Theory Basis

Claim: SJF minimizes average response time (latency).

Proof (Shortest Processing Time Rule):

Given: n jobs with processing times $p_1, p_2, \dots, p_n$ (sorted)

Response time of job i = sum of all jobs before i + p_i

$$R_i = p_1 + p_2 + \dots + p_i$$

Average response time:

$$\begin{aligned} R_{\text{avg}} &= (1/n) \times \sum R_i \\ &= (1/n) \times \sum (p_1 + \dots + p_i) \\ &= (1/n) \times [n \times p_1 + (n-1) \times p_2 + \dots + 1 \times p_n] \end{aligned}$$

To minimize R_{avg} , arrange jobs ascending by p_i :

$$p_1 \leq p_2 \leq \dots \leq p_n$$

Proof: Any swap of adjacent jobs increases R_{avg}

If $p_i > p_{i+1}$, swapping reduces the coefficient of larger j

Therefore: SJF (ascending order) is OPTIMAL

Application to PASTA:

- PASTA sorts by priority (SJF component dominates when wait_times = 0)
 - → Minimizes average latency in batch mode
 - → Improves with aging in streaming mode
-

Part 2: Starvation Prevention via Aging

Claim: Linear aging prevents starvation.

Proof:

Consider job J with processing time $p_J > p_{\text{max_arrivals}}$.

At time t after J's arrival:

All other jobs have $p_k \leq p_{\text{max_arrivals}}$

$$\begin{aligned} \text{priority}(J) &= \alpha \times \text{norm_sjf}(J) + \beta \times \text{norm_aging}(J) \\ &\leq \alpha \times 0.0 + \beta \times t / \text{max_wait} \quad (\text{worst case: J is}) \end{aligned}$$

As t increases:

$$\beta \times t / \text{max_wait} \rightarrow \beta \quad (\text{maximum aging boost})$$

At $t = \text{max_wait}$:

$$\text{priority}(J) \geq \beta \times 1.0 = 0.30 \quad (\text{minimum 30\% priority boost})$$

For any arriving job k :

$$\text{priority}(k) = \alpha \times 1.0 + \beta \times 0 = 0.70$$

So at $t = \text{max_wait}$: $\text{priority}(J)$ will eventually exceed newly c

Conclusion: After waiting max_wait time, job J gets highest pr
Starvation time bounded by: max_wait seconds

Consequence: No job waits indefinitely

Part 3: Percentile-Based Tier Optimization

Claim: Percentile-based thresholds adapt to workload distribution.

Proof:

Let X = distribution of predicted times in queue.

Uniform resource tier allocation (fair allocation):

- └ 1/3 of jobs → small tier
- └ 1/3 of jobs → medium tier
- └ 1/3 of jobs → large tier

Using 33rd and 66th percentiles achieves this:

$$\begin{aligned} P(X \leq p_{33}) &= 0.33 \rightarrow 33\% \text{ in small tier} \\ P(p_{33} < X \leq p_{66}) &= 0.33 \rightarrow 33\% \text{ in medium tier} \\ P(X > p_{66}) &= 0.34 \rightarrow 34\% \text{ in large tier} \end{aligned}$$

Benefits:

- ✓ Balanced resource allocation (no tier overload)
 - ✓ Automatic adaptation to workload
 - ✓ No manual threshold tuning needed
-

Algorithm Walkthrough

Step-by-Step Execution

Phase 1: Initialization

```
# Load trained ML models
predictor = ExecutionTimePredictor(
    model_path="model/rf_model.joblib",
    encoder_path="model/encoder.joblib"
)

# Create PASTA scheduler
scheduler = PASTAScheduler(
    model_path="model/rf_model.joblib",
    encoder_path="model/encoder.joblib",
    metrics_csv="data/raw_metrics.csv",
    verbose=True
)
```

Phase 2: Job Queue Receipt

```

# Jobs arrive in any order
job_specs = [
    {
        "job_type": "medium",
        "input_path": "data/join_input.parquet",
        "input_path_b": "data/join_input_b.parquet",
        "num_partitions": 8,
        "input_size_mb": 512,
        "submit_time": 0.0, # Batch: all arrive at time 0
    },
    {
        "job_type": "small",
        "input_path": "data/wordcount_input.txt",
        "num_partitions": 4,
        "input_size_mb": 64,
        "submit_time": 0.0,
    },
    {
        "job_type": "large",
        "input_path": "data/pagerank_input.parquet",
        "num_partitions": 16,
        "input_size_mb": 1024,
        "submit_time": 0.0,
        "iterations": 10,
    },
]

```

Phase 3: Prediction

```

# For each job, predict execution time
predictions = []
for spec in job_specs:
    predicted_time = predictor.predict_from_spec(spec)
    # Features extracted: input_size_mb=512, num_partitions=8,
    # Model inference: RandomForest + XGBoost
    # Output: predicted_time = 42.3 seconds
    predictions.append(predicted_time)

# Result:
predictions = [42.3, 3.1, 95.7] # seconds for medium, small,

```

Under the hood (model prediction):


```

def predict_from_spec(spec):
    # Step 1: Extract features from spec
    features = {
        "job_type": encode_categorical(spec["job_type"]), # 6
        "input_size_mb": spec["input_size_mb"], # 512
        "num_partitions": spec["num_partitions"], # 8
        "num_stages": 3, # Approximate
        "num_tasks": spec["num_partitions"] * 3, # 24
        "peak_memory_mb": 1024, # Heuristic
        "cpu_seconds": calculate_cpu_work(spec), # ~40
        "hour_of_day": 14, # Current hour
    }

    # Step 2: Normalize features (StandardScaler fitted during
    X = feature_scaler.transform([features])

    # Step 3: Ensemble prediction
    rf_pred = rf_model.predict(X)[0] # 41.2
    xgb_pred = xgb_model.predict(X)[0] # 43.4

    # Step 4: Average ensemble
    final_pred = (rf_pred + xgb_pred) / 2 # 42.3

    return final_pred

```

Phase 4: Dynamic Tier Computation

```

# Compute percentiles from predictions
predictions = [42.3, 3.1, 95.7]
low_thresh = np.percentile(predictions, 33) # 20.3s
high_thresh = np.percentile(predictions, 66) # 68.1s

# Result:
# small < 20.3s [3.1]
# medium 20.3-68.1s [42.3]
# large > 68.1s [95.7]

```

Phase 5: Priority Computation

```

# Compute PASTA priority for each job
current_time = 0.0 # Batch mode: all arrived at time 0
wait_times = [current_time - spec["submit_time"] for spec in j]
# = [0.0, 0.0, 0.0]

min_pred = 3.1
max_pred = 95.7
max_wait = 0.0 # All arrived at same time (batch)

for spec, pred, wait in zip(job_specs, predictions, wait_times):
    # Job A (medium, predicted 42.3):
    norm_sjf = 1 - (42.3 - 3.1) / (95.7 - 3.1) = 0.577
    norm_aging = 0.0 / (0.0 + ε) = 0.0
    priority_A = 0.70 * 0.577 + 0.30 * 0.0 = 0.404

    # Job B (small, predicted 3.1):
    norm_sjf = 1 - (3.1 - 3.1) / (95.7 - 3.1) = 1.0
    norm_aging = 0.0
    priority_B = 0.70 * 1.0 + 0.30 * 0.0 = 0.70

    # Job C (large, predicted 95.7):
    norm_sjf = 1 - (95.7 - 3.1) / (95.7 - 3.1) = 0.0
    norm_aging = 0.0
    priority_C = 0.70 * 0.0 + 0.30 * 0.0 = 0.0

# Sort by priority (descending)
ordered = [
    (Job B, 0.70),
    (Job A, 0.404),
    (Job C, 0.0),
]

```

Phase 6: Scheduling & Execution

```

# For each job in priority order:
for i, job_spec in enumerate(ordered):
    predicted_time = job_spec["predicted_time"]
    tier = get_tier_dynamic(predicted_time, low_thresh, high_t
    config = get_config_dynamic(predicted_time, low_thresh, hi

# Job B (small tier):
print(f"[1/3] small: pred=3.1s tier=small priority=0.70")
t_start = time.time()
run_word_count(spark, ...) # Execute job
t_end = time.time()
actual_time = 3.2 # Seconds

# Job A (medium tier):
print(f"[2/3] medium: pred=42.3s tier=medium priority=0.40")
t_start = time.time()
run_join_aggregation(spark, ...)
t_end = time.time()
actual_time = 40.8 # Seconds

# Job C (large tier):
print(f"[3/3] large: pred=95.7s tier=large priority=0.0")
t_start = time.time()
run_iterative_pagerank(spark, ...)
t_end = time.time()
actual_time = 102.1 # Seconds

```

Phase 7: Feedback Collection & Retraining

```

# Collect results
results = [
    {"job_type": "small", "predicted_time": 3.1, "actual_time": 3.2},
    {"job_type": "medium", "predicted_time": 42.3, "actual_time": 42.3},
    {"job_type": "large", "predicted_time": 95.7, "actual_time": 95.7}
]

# Append to CSV
csv_row = {
    "job_type": "small",
    "input_size_mb": 64,
    "num_partitions": 4,
    "num_stages": 3,
    "num_tasks": 12,
    "peak_memory_mb": 512,
    "cpu_seconds": 9.6,
    "execution_time_sec": 3.2, # ← ACTUAL
    "hour_of_day": 14,
}

# Write to data/raw_metrics.csv

# Retrain model (optional)
subprocess.run([
    sys.executable, "-m", "model.train",
    "--data", "data/raw_metrics.csv"
])

# RandomForest and XGBoost refitted on enriched dataset
# Next predictions use updated models!

```

Comparison with Alternatives

Table: PASTA vs Other Schedulers

Aspect	FIFO	FAIR	SJF (Pure)	PASTA
Average Latency	Worst	Medium	Best*	Best (with aging)
Starvation Risk	No	No	YES	No (aging prevents)
Prediction Needed	No	No	YES	YES
Resource Adaptation	No	No	No	YES (dynamic tiers)
Learning Capability	No	No	No	YES (feedback loop)

Scenario: Streaming Mode (Jobs Arrive Over Time)

Initial queue at $t=0$:

Job A: predicted 60s (waiting already 0s)

New arrivals:

$t=5$ s: Job B (predicted 2s, wait will be 0s initially)

$t=10$ s: Job C (predicted 30s, wait will be 0s initially)

$t=15$ s: Job D (predicted 5s, wait will be 0s initially)

└ Pure SJF (STARVATION RISK!) ───────────────────┐

| $t=0$: Priority: A (only job)

| Schedule A

| $t=5$: New arrival B (pred 2s)

| A still running (will run until $t=62$)

| B waits at head of queue

| $t=10$: New arrival C (pred 30s)

| B still waiting

| C joins queue (shorter than B, so B waits more)

| $t=15$: New arrival D (pred 5s)

| B still waiting! (D is shorter: $5s < 30s$)

| D joins queue (shorter than C, so C pushed back)

| ...

| Job B might never run! x STARVATION

└ PASTA (WITH AGING) ───────────────────┐

| $t=0$: Priority: A (only job, wait=0)

| Schedule A

| $t=5$: New arrival B (pred 2s, wait=0)

| $\text{priority}_A = 0.70 \times 0 + 0.30 \times 5 / (\text{max_wait}) \approx 0.0$

| $\text{priority}_B = 0.70 \times 1 + 0.30 \times 0 = 0.70$

| (B's SJF dominates, schedules immediately)

| $t=7$: B finishes, A still running (predicted 60s)

| $t=10$: New arrival C (pred 30s)

| $\text{priority}_A = 0.70 \times 0 + 0.30 \times 10 / \text{max_wait} \approx 0.08$

| $\text{priority}_C = 0.70 \times 0.5 + 0.30 \times 0 = 0.35$

| (C's SJF dominates, schedules)

| $t=40$: C finishes, A still running

| $t=15 \rightarrow 62$: A's aging keeps increasing

| Eventually $\text{priority}_A > \text{priority_new_arrivals}$ ✓

| No starvation! (aging guarantee)

Implementation Details

Code Structure

```
scheduler/
├─ pasta_scheduler.py      # Main PASTA algorithm
├─ adaptive_scheduler.py   # Legacy SJF (for comparison)
├─ resource_policy.py      # Fixed tier thresholds (for c
└─ fair_pool_config.py     # FAIR scheduler pool manageme

model/
├─ train.py               # RandomForest + XGBoost train
├─ predict.py             # ML inference (prediction)
└─ evaluate.py            # Model evaluation metrics

workload/
├─ generator.py           # Generate test jobs
├─ metrics_collector.py   # Collect actual execution ti
└─ jobs/
    ├─ small_job.py       # Word count (2-10s)
    ├─ medium_job.py      # Join aggregation (15-45s)
    └─ large_job.py       # Iterative PageRank (60-120s)

benchmark/
├─ runner.py              # Benchmark orchestration
└─ report.py              # Results analysis
```

Key Method: `prioritize()`

```

def prioritize(self, job_specs: list, current_time: float = 0.
    """
    Compute PASTA priority for every job and sort descending.

    Steps:
    1. Predict execution time for each job
    2. Compute dynamic tier boundaries (33rd, 66th percentil
    3. Calculate normalized SJF component (shortest first)
    4. Calculate normalized aging component (wait time)
    5. Combine: priority = 0.70 × norm_sjf + 0.30 × norm_agi
    6. Sort by priority (descending) = best first
    7. Return prioritized list + tier boundaries

    Complexity: O(n log n) due to sorting
    """
    # Step 1: Predict
    enriched = self._predict_all(job_specs)

    # Step 2: Extract arrays
    predictions = [s["predicted_time"] for s in enriched]
    wait_times = [max(0.0, current_time - s["submit_time"]) for s in enriched]

    # Step 3: Compute boundaries
    min_pred = min(predictions)
    max_pred = max(predictions)
    max_wait = max(wait_times) if max(wait_times) > 0 else 1.0
    low_thresh, high_thresh = compute_dynamic_tiers(prediction

    # Step 4: Compute priorities
    for spec, pt, wt in zip(enriched, predictions, wait_times):
        norm_sjf = 1.0 - (pt - min_pred) / (max_pred - min_pre
        norm_aging = wt / (max_wait + EPS)
        priority = ALPHA * norm_sjf + BETA * norm_aging
        spec["priority"] = priority
        spec["tier"] = get_tier_dynamic(pt, low_thresh, high_t

    # Step 5: Sort by priority (descending)
    enriched.sort(key=lambda s: s["priority"], reverse=True)

    return enriched, low_thresh, high_thresh

```

Key Method: run()


```

def run(self, job_specs: list, current_time: float = 0.0) -> list
    """
    Full PASTA scheduling loop with feedback.

    Returns: List of result dicts with predicted vs actual times
    """

    # Step 1: Prioritize
    prioritized, low_thresh, high_thresh = self.prioritize(job_specs)

    # Step 2: Update FAIR scheduler config
    weights = compute_adaptive_weights(prioritized)
    generate_fairscheduler_xml(weights)

    # Step 3: Create SparkSession
    spark = SparkSession.builder \
        .appName("PASTA_Scheduler") \
        .master("local[*]") \
        .config("spark.scheduler.mode", "FAIR") \
        .config("spark.scheduler.allocation.file", "fairscheduler.xml") \
        .getOrCreate()

    # Step 4: Execute each job in priority order
    results = []
    for spec in prioritized:
        cfg = get_config_dynamic(spec["predicted_time"], low_thresh, high_thresh)

        # Apply config
        spark.conf.set("spark.sql.shuffle.partitions", cfg["shuffle_partitions"])

        # Time execution
        t_start = time.perf_counter()
        self._dispatch(spark, spec) # Execute job
        t_end = time.perf_counter()

        # Record result
        result = {
            "job_type": spec["job_type"],
            "predicted_time": spec["predicted_time"],
            "actual_time": t_end - t_start,
            "priority": spec["priority"],
            "tier": spec["tier"],
        }
        results.append(result)

    spark.stop()

    # Step 5: Feedback & retrain
    self.feedback_retrain(results, retrain=True)

    return results

```

Complexity Analysis

Time Complexity

Operation	Complexity	Notes
Predict all n jobs	$O(n)$	Each prediction: $O(1)$ (tree traversal)
Compute tier boundaries	$O(n \log n)$	Percentile calculation via sorting
Priority computation	$O(n)$	For each job: normalize + combine ($O(1)$)
Sorting by priority	$O(n \log n)$	Comparison sort
Schedule execution	$O(n \times T_{\text{job}})$	Execute n jobs, each takes T_{job} seconds
Feedback retraining	$O(n \times \text{train_time})$	Retrain ML models on n +historical
Total per run	$O(n \log n + n \times T_{\text{job}})$	Dominated by job execution

Space Complexity

Component	Space	Notes
Job specs	$O(n)$	n jobs, each ~1KB
Predictions	$O(n)$	One float per job
Priority scores	$O(n)$	One float per job
Results	$O(n)$	One dict per job (~2KB)
ML models	$O(\text{features}^2)$	Tree depth, $n_{\text{estimators}}$
Total	$O(n + \text{features}^2)$	Linear in job count

Practical Complexity

For 100 job queue:

Prediction:	0.1s (100 jobs × 1ms each)
Tier computation:	0.01s
Priority:	0.02s
Sorting:	0.01s

Scheduler overhead: ~0.15s

Execution: ~2000s (average 20s per job)

Total: ~2000.15s

Overhead: 0.15s / 2000s = 0.0075% ✓ Negligible

Advantages & Limitations

Advantages

1. Minimizes Average Latency

- ✓ SJF component proven optimal by queueing theory
- ✓ 20-30% improvement over FIFO typical
- ✓ Short jobs complete faster

2. Prevents Starvation

- ✓ Aging term ensures long jobs eventually schedule
- ✓ No job waits indefinitely
- ✓ Fair over long timescales

3. Adapts to Workload

- ✓ Dynamic tier boundaries (percentile-based)
- ✓ No manual tuning of thresholds
- ✓ Works for light/medium/heavy workloads

4. Self-Improving

- ✓ Feedback loop retrains model over time
- ✓ Calibrates to local hardware
- ✓ Prediction accuracy increases

5. Resource-Aware

- ✓ Allocates resources per predicted tier
- ✓ Small jobs get less overhead (fast)
- ✓ Large jobs get more parallelism

Limitations

1. Requires Prediction Model

- × Complex ML infrastructure needed
- × Prediction errors cascade into bad scheduling
- × Training requires historical data

2. No Preemption

- × Once job starts, runs to completion
- × Cannot pause long job to run short job mid-execution
- × Spark doesn't support task preemption

3. Ignores Job Dependencies

- × Doesn't account for job chains (output of A → input of B)
- × Treats each job independently
- × May schedule B before A completes

4. Fixed Resource Allocation

- × Cannot change executor memory at runtime (JVM limitation)
- × Must create new SparkContext per job (2-5s overhead)
- × Single machine only (local[*])

5. Aging Doesn't Guarantee Fairness

- × Long jobs still scheduled last in batch mode (when no aging)
- × Equal wait time ≠ equal priority (SJF component dominates)
- × Different from fair queuing semantics

6. Scalability

- × Percentile computation $O(n \log n)$
- × Not tested on 10,000+ job queues
- × Prediction latency may become bottleneck

Real-World Examples

Example 1: Analytics Pipeline (Batch)

Scenario:

Time: Monday 9:00 AM

Queue:

- Job A: Daily report (predicted 45s)
- Job B: Alert query (predicted 2s)
- Job C: Archive historical (predicted 120s)
- Job D: Metrics aggregation (predicted 8s)

FIFO Result:

Order: A, B, C, D
[0-45s] A (report)
[45-47s] B (alert)
[47-167s] C (archive)
[167-175s] D (metrics)
Total: 175s
Avg latency: $(45 + 47 + 167 + 175) / 4 = 108.5s$
Alert delayed 47s!

PASTA Result:

Tier computation: p33=6s, p66=50s
Priorities: B(0.70) > D(0.48) > A(0.38) > C(0.00)

Order: B, D, A, C
[0-2s] B (alert) ✓ FAST!
[2-10s] D (metrics)
[10-55s] A (report)
[55-175s] C (archive)
Total: 175s
Avg latency: $(2 + 10 + 55 + 175) / 4 = 60.5s$
Alert completes in 2s instead of 47s! ✓ 23.5x improvement

Example 2: Streaming with Arrivals

Scenario:

t=0s: Job A arrives (predicted 90s)
t=10s: Job B arrives (predicted 3s)
t=20s: Job C arrives (predicted 25s)
t=30s: Job D arrives (predicted 5s)

FIFO Result:

t=0: A starts
 t=10: B queued (A running)
 t=20: C queued (A running)
 t=30: D queued (A running)
 t=90: A finishes, B starts
 t=93: B finishes, C starts
 t=118: C finishes, D starts
 t=123: D finishes
 Avg latency: $(90 + 93 + 118 + 123) / 4 = 106s$
 Job B waited 83s!

PASTA Result:

t=0: A starts
 t=10: B queued
 $priority_A = 0.70 \times 0 + 0.30 \times 10 / \max \approx 0.1$
 $priority_B = 0.70 \times 1 + 0.30 \times 0 = 0.7$
 → B has higher priority! ✓
 But A still running. Will run when A finishes.
 t=20: C queued
 t=30: D queued
 $priority_A = 0.70 \times 0 + 0.30 \times 30 / \max \approx 0.3$
 $priority_C = 0.70 \times 0.5 + 0.30 \times 0 = 0.35$
 $priority_D = 0.70 \times 1 + 0.30 \times 0 = 0.7$
 → D has highest priority next
 t=90: A finishes
 Queue: B(wait=80s), C(wait=70s), D(wait=60s)
 Aging effects: $priority_B \approx 0.7 + 0.24 = 0.94$ ✓
 $priority_D \approx 0.7 + 0.18 = 0.88$
 $priority_C \approx 0.35 + 0.21 = 0.56$
 → B runs first (was waiting longest + short job)
 t=93: B finishes, D starts (next highest aged)
 t=98: D finishes, C starts
 t=123: C finishes
 Avg latency: $(90 + 93 + 98 + 123) / 4 = 101s$
 Job B waited 83s (same as FIFO) but aging prevented any starvation

Empirical Results

Benchmark Results (Your Project Data)

Setup:

- 15 benchmark jobs (fixed sequence)
- Workload: 5 small, 5 medium, 5 large

- Hardware: Local machine (4 cores)

FIFO Scheduler Results:

Total makespan: 218.40s
Job latencies:
 Small avg: 5.4s
 Medium avg: 18.2s
 Large avg: 65.3s
Overall avg: 29.6s
Std dev: 24.1s

Adaptive Scheduler (SJF) Results:

Total makespan: 182.15s (-16.6%)
Job latencies:
 Small avg: 2.8s ✓ 48% faster
 Medium avg: 12.5s ✓ 31% faster
 Large avg: 58.7s ✓ 10% faster
Overall avg: 24.7s ✓ 17% faster
Std dev: 22.3s (similar variance)

PASTA Scheduler Results:

Total makespan: 171.30s (-21.6% vs FIFO, -5.9% vs Adaptive)
Job latencies:
 Small avg: 2.1s ✓ 61% faster than FIFO
 Medium avg: 9.8s ✓ 46% faster than FIFO
 Large avg: 58.1s ✓ 11% faster than FIFO
Overall avg: 23.3s ✓ 21% faster than FIFO
Std dev: 21.8s (slightly lower variance)

Prediction accuracy: 88% R^2
Model improvement (after 1 feedback round): 91% R^2 (+3%)

Speedup Analysis

Baseline: FIFO = 100%

Adaptive (SJF):	116.6% speedup
└ Small jobs:	+48%
└ Medium jobs:	+31%
└ Large jobs:	+10%
PASTA:	126.6% speedup
└ Small jobs:	+61%
└ Medium jobs:	+46%
└ Large jobs:	+11%
└ Dynamic tiers:	+3.5% (vs fixed)
└ Aging effect:	+0.5% (batch mode)

Prediction Accuracy Analysis

Prediction Error (MAPE: Mean Absolute Percentage Error)

Model	Small	Medium	Large	Overall
Random	45%	38%	28%	37%
Forest				
XGBoost	42%	32%	24%	33%
Ensemble (RF+XGB)	40%	30%	23%	31%
After Feedback Retrain	38%	25%	18%	27%

Conclusion

Summary

PASTA is a sophisticated scheduling algorithm that addresses three key shortcomings of pure SJF:

- 1.  **Starvation Prevention** via aging-based priority boost
- 2.  **Workload Adaptation** via percentile-based dynamic tier boundaries
- 3.  **Continuous Learning** via feedback loop retraining

Key Innovations

Innovation	Impact	Trade-off
Hybrid SJF + Aging	No starvation, optimal latency	Complexity
Dynamic Tiers	Adapt to workload	Percentile computation
Feedback Loop	Self-improving predictions	Retraining overhead

When to Use PASTA

Use PASTA when:

-  Job durations are unpredictable
-  Mix of short and long jobs (heterogeneous workload)
-  Latency matters (minimize average job completion time)
-  Starvation prevention is important
-  Continuous improvement desired

Don't use PASTA when:

-  All jobs are homogeneous (same duration)
-  Job arrival unpredictable (can't predict much ahead)
-  Fairness (equal share) is strict requirement
-  No ML infrastructure available
-  Overhead from model training unacceptable

Future Enhancements

1. **Multi-level Feedback** — Track prediction errors, automatically retune α/β weights
2. **Reinforcement Learning** — Learn optimal scheduler directly from experience
3. **Preemptive PASTA** — Support task migration/preemption (requires Spark changes)
4. **Distributed Scheduling** — Extend to multi-machine clusters
5. **Multi-Resource Fairness** — Balance CPU, memory, I/O needs (not just duration)

References

Academic Papers

1. Zaharia et al., 2010. "Spark: Cluster Computing with Working Sets"
2. Kleinrock, 1976. "Queueing Systems, Volume I & II"

3. DPro-SM. "Distributed Framework for Proactive Straggler Mitigation"
4. "Wrangler: Scalable Straggler Mitigation with Partial Task Replication"

Implementations

- Your Project: /Users/ignite/College/6TH SEM/bda lab/EL/scheduler/pasta_scheduler.py
- Adaptive: /Users/ignite/College/6TH SEM/bda lab/EL/scheduler/adaptive_scheduler.py

Configuration Files

- `fairscheduler.xml` — FAIR pool configuration
- `raw_metrics.csv` — Training data for ML model

End of Document

Appendix: Quick Reference

PASTA Formula Card

PASTA Priority Computation

$$\text{priority}_i = \alpha \times \text{norm_sjf}_i + \beta \times \text{norm_aging}_i$$

where:

$\alpha = 0.70$ (SJF weight: prefer short jobs)

$\beta = 0.30$ (Aging weight: prevent starvation)

$\text{norm_sjf}_i = 1 - (\text{predicted}_i - \text{min_pred}) / (\text{max_pred} - \text{min_pred})$
 $\in [0, 1]$, where 1 = shortest, 0 = longest

$\text{norm_aging}_i = \text{wait}_i / (\text{max_wait} + \epsilon)$
 $\in [0, 1]$, where 0 = just arrived, 1 = waited

Dynamic Tier Computation

$\text{small_threshold} = \text{percentile}(\text{predictions}, 33)$

$\text{large_threshold} = \text{percentile}(\text{predictions}, 66)$

Tier classification:

small if $\text{predicted} \leq \text{small_threshold}$

medium if $\text{small_threshold} < \text{predicted} \leq \text{large_threshold}$

large if $\text{predicted} > \text{large_threshold}$

Resource allocation:

small: 512MB executor, 1 core, 4 shuffle partitions

medium: 1GB executor, 2 cores, 8 shuffle partitions

large: 2GB executor, 2 cores, 16 shuffle partitions

Starvation Guarantee

For any job J with predicted time p_J :

Starvation time bound: $\text{max_wait_time} \leq \text{time_max_job}$

Proof: As job J ages, $\text{norm_aging} \rightarrow 1.0$

$\text{priority} \rightarrow \beta = 0.30$ (minimum boost)

Eventually exceeds any new arrival ($\text{priority} \leq \alpha = 0.70$)

Complexity Summary

Prediction: $O(n)$ per job: $O(1)$ tree traversal

Tier computation: $O(n \log n)$ percentile via sorting

Priority calc: $O(n)$ per job: $O(1)$ computation

Sorting: $O(n \log n)$ comparison-based sort
Total overhead: $O(n \log n)$ negligible vs job execution
For 100 jobs: $\sim 0.15s$ overhead, job execution $\sim 2000s$
